

GoldenMinutes

AI Hackathon 2026 — Judge Feedback, One Criterion at a Time

Overall Score	74.67 / 100
Criteria	7

The following pages organize the feedback criterion-by-criterion, preserving the three judges' comments shown in the supplied screenshots.

1. Problem Relevance

Score: 17.00 / 20

Judge	Feedback
Judge 1	Focuses on stopping scam payments before cash out. Check the response time needed by fraud teams.
Judge 2	Addresses a highly relevant MFS fraud problem by focusing on real-time intervention before suspicious funds are cashed out, although the real-world problem claims still need stronger primary-source validation.
Judge 3	Richt zich zeer effectief op de kritieke tijdsdruk ("golden window") bij MFS-fraude in Bangladesh, waarbij het geld snel via katvangers wordt opgenomen.

Criterion performance: 85.0% of available points.

2. AI/ML Depth

Score: 14.67 / 20

Judge	Feedback
Judge 1	Uses payment patterns and wallet links to find risk. Show whether wallet links improve detection.
Judge 2	Strong combination of LightGBM, Isolation Forest, graph-based intelligence, SHAP, rules, risk fusion, and held-out fraud-type testing, though advanced graph learning/GNN is not actually implemented.
Judge 3	Indrukwekkend gelaagde AI-architectuur die LightGBM-risicoscoring, Isolation Forest-anomaliedetectie, NetworkX-graafeigenschappen en SHAP-verklaarbaarheid combineert met ablatiestudies.

Criterion performance: 73.4% of available points.

3. Business / Customer Impact

Score: 15.33 / 20

Judge	Feedback
Judge 1	Could reduce fraud losses, but claimed savings need testing. Measure losses prevented and unnecessary payment delays.
Judge 2	Has clear potential to reduce fraud losses while controlling false customer friction, but the impact has not yet been demonstrated with real MFS transactions or customers.
Judge 3	Biedt grote meerwaarde door wrijving te personaliseren (Allow, Warn, Verify, Hold) en transacties te prioriteren op basis van risicobedrag en resterende tijd tot cash-out.

Criterion performance: 76.6% of available points.

4. Prototype Quality

Score: 10.33 / 15

Judge	Feedback
Judge 1	A local demo is documented, but no public link was provided. Share a working demo link and fix model loading.
Judge 2	Provides a fairly complete end-to-end prototype covering scoring, graph analysis, customer warnings, analyst review, APIs, and feedback, but it remains a synthetic-data prototype rather than a live MFS deployment.
Judge 3	Uitstekend uitgewerkt prototype met FastAPI-backend, React UI, en analytische console, waarschuwingstermes voor klanten in het Bengaals/Engels en scriptbare aanvalssimulaties.

Criterion performance: 68.9% of available points.

5. Innovation

Score: 7.67 / 10

Judge	Feedback
Judge 1	Combines wallet checks with urgent human review. Test scams that are harder to spot.
Judge 2	The combination of golden-window prioritization, recipient-network intelligence, proportionate intervention, Bangla explanations, and human review provides good practical innovation, although most underlying AI techniques are established.
Judge 3	Vernieuwende benadering door de "golden window"-prioritering, netwerkindelligentie aan de ontvangende kant en gefundeerde LLM-samenvatting te combineren.

Criterion performance: 76.7% of available points.

6. Scalability & Integration

Score: 6.67 / 10

Judge	Feedback
Judge 1	The design needs shared data for larger use. Test several scoring servers using the same data.
Judge 2	The API-based architecture provides a reasonable path toward MFS integration, but large-scale performance and integration with an actual MFS infrastructure have not been demonstrated.
Judge 3	Goed gestructureerde modulaire API-architectuur met SQLite/SQLAlchemy backend, hoewel graafberekeningen in productie moeten worden opgeschaald naar dedicated graaf-databases.

Criterion performance: 66.7% of available points.

7. Responsible AI & Security

Score: 3.00 / 5

Judge	Feedback
Judge 1	Human review helps, but shared access keys weaken protection. Give each reviewer a secure personal login.
Judge 2	Strong consideration of human review, explainability, fairness, auditability, limited LLM authority, and prompt-injection risks, although these safeguards have not yet been validated in production.
Judge 3	Uitstekende waarborgen voor verantwoordelijke AI: menselijke controle bij ingehouden transacties, LLM beperkt tot feitelijke redenodes, bestand tegen prompt-injectie en demografische eerlijkheidscontroles.

Criterion performance: 60.0% of available points.